

# Deep Learning

# The Revolution

Deep learning has provided **empirically much better** methods for:


- Hard prediction problems
- Generative models of natural data distributions

Especially over high-dimensional data such as images, video, speech, text, and robotic control

# Deep Learning for Speech Recognition

- The first breakthrough results of “deep learning” on large datasets happened in speech recognition
- Context-Dependent Pre-trained Deep Neural Networks for Large Vocabulary Speech Recognition. Dahl et al. (2010)

Acoustic model Measuring WER	RT03S FSH	Hub5 SWB
Traditional (GMM) (2012)	27.4	23.6
Deep Learning (Dahl et al. 2012)	18.5 (-33%)	16.1 (-32%)
Xiong et al. (2017)		5.8





[www.image-net.org](http://www.image-net.org)

**22K** categories and **14M** images

- Animals
  - Bird
  - Fish
  - Mammal
  - Invertebrate
- Plants
  - Tree
  - Flower
  - Food
  - Materials
- Structures
  - Artifact
    - Tools
    - Appliances
    - Structures
- Person
  - Scenes
    - Indoor
    - Geological Formations
  - Sport Activities



# IMAGENET Large Scale Visual Recognition Challenge

The Image Classification Challenge:  
1,000 object classes  
1,431,167 images



Output:  
Scale  
T-shirt  
Steel drum  
Drumstick  
Mud turtle




Output:  
Scale  
T-shirt  
Giant panda  
Drumstick  
Mud turtle





# IMAGENET Large Scale Visual Recognition Challenge


The Image Classification Challenge:  
1,000 object classes  
1,431,167 images



# Reinforcement Learning (AlphaGo)

The model works from a 319-dimensional input representing the board and uses a regression model to score potential next moves

Combined with Monte Carlo Tree Search, this “solved” Go much more quickly than anyone had been imagining



# Recipe for winning Kaggle competitions


1. Careful data preprocessing, cleaning, augmentation, and feature engineering (this hasn't gone away to win Kaggle!)
2.
  - a. **For classic, structured data tables: Gradient-boosted decision trees** (xgboost). Roughly, improved MART.
  - b. **For “unstructured” text, images, video, speech: Neural networks**
3. Ensembling/stacking of models, with careful cross-validation testing to find best final configuration



# What is deep learning (DL)?


- **Deep learning** is a subfield of **machine learning (statistics?)**
- Most machine learning methods work well because of **human-designed input features or representations**
  - SIFT or HOG features for vision
  - MFCC or LPC features for speech
  - Features about words parts (suffix, capitalized?) for finding person or location names
- Machine learning becomes just optimizing weights to best make a final prediction

Feature	NER
Current Word	✓
Previous Word	✓
Next Word	✓
Current Word Character n-gram	all
Current POS Tag	✓
Surrounding POS Tag Sequence	✓
Current Word Shape	✓
Surrounding Word Shape Sequence	✓
Presence of Word in Left Window	size 4
Presence of Word in Right Window	size 4




# What is deep learning (DL)?


- In contrast to standard machine learning,
- Representation learning attempts to automatically learn good features or representations
- Deep learning algorithms learn multiple levels of representations (here:  $h^1, h^2, h^3$ ) and an output ( $h^4$ )
- From “raw” inputs  $\mathbf{x}$  (e.g. sound, pixels, characters, or words)
- Neural networks are the currently successful method for deep learning
- A.k.a. “Differentiable Programming”



# A Brief History of Neural Nets




S. McCulloch - W. Pitts



- Adjustable Weights
- Weights are not Learned



F. Rosenblatt




- Learnable Weights and Threshold



B. Widrow - M. Hoff




M. Minsky - S. Papert



- XOR Problem




D. Rumelhart - G. Hinton - R. Williams



- Solution to nonlinearly separable problems
- Big computation, local optima and overfitting




V. Vapnik - C. Cortes



- Limitations of learning prior knowledge
- Kernel function: Human Intervention




G. Hinton - S. Ruslan



- Hierarchical feature Learning



# Algorithms




# Data



# Computation




# Neurons in the brain



- Axons
- Dendrites
- Synapses
- Receptors

## A single neuron

A computational unit with  $n$  (3) inputs and 1 output and parameters  $w, b$







# A neuron is essentially a logistic regression unit

$f$  = nonlinear activation function (e.g., logistic),  
 $w$  = weights,  $b$  = bias,  $h$  = hidden,  $x$  = inputs

$$h_{w,b}(x) = f(w^T x + b)$$


$b$ : We can have an “always on” feature, which gives a class prior, or separate it out, as a bias term

$$f(z) = \frac{1}{1 + e^{-z}}$$

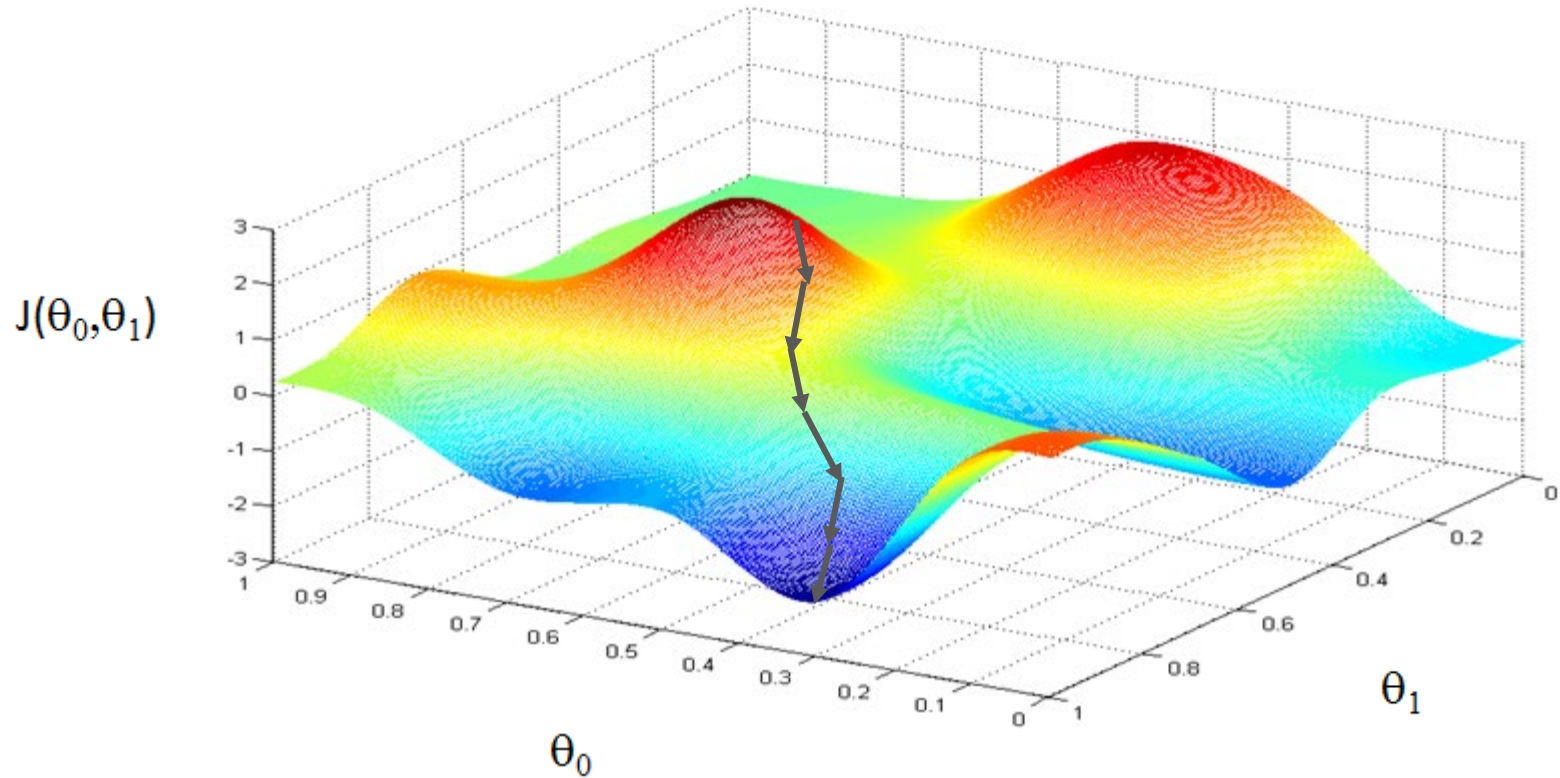


$w, b$  are the parameters of this neuron  
i.e., this logistic regression model

A neural net = running many logistic regressions



Use gradient descent to learn weights and biases





# Gradient descent

$$\begin{array}{l} \text{repeat until convergence } \{ \\ \quad \theta_j := \theta_j - \alpha \frac{\partial}{\partial \theta_j} J(\theta_0, \theta_1) \\ \} \end{array}$$



- Calculate the gradient of the loss function w.r.t. parameters
- Determine the learning rate
- Instead of batch gradient descent, we mostly use **stochastic gradient descent (SGD)**, which is much **faster**, easier to escape **local minima** and more stable.

# Backpropagation algorithm

**Input:** Training set  $\{(\mathbf{x}^{(i)}, \mathbf{y}^{(i)})\}_{i=1}^m$

1. Set  $\Delta^{(l)} = 0$  for all  $l$
  2. **for**  $i = 1$  **to**  $m$
  3.     Set  $\mathbf{a}^{(0)} = \mathbf{x}^{(i)}$
  4.     Perform forward propagation to calculate  $\mathbf{a}^{(l)}$  for  $l = 1, \dots, L$
  5.     Using  $y^{(i)}$ , compute  $\delta^{(L)} = \mathbf{a}^{(L)} - \mathbf{y}^{(i)}$
  6.     Compute  $\delta^{(L-1)}, \dots, \delta^{(1)}$
  7.      $\Delta^{(l)} := \Delta^{(l)} + \delta^{(l)}(\mathbf{a}^{(l-1)})^T$  for all  $l$
  8.      $\Delta^{(l)} := \frac{1}{m}\Delta^{(l)}$  for all  $l$
  9.      $W^{(l)} := W^{(l)} - \eta\Delta^{(l)}$  for all  $l$
-

# Automatic Differentiation



- The gradient computation can be automatically inferred from the symbolic expression of the fprop.
- Each node type needs to know how to compute its output and how to compute the gradient wrt its inputs given the gradient wrt its output.
- Modern DL frameworks (Tensorflow, PyTorch, etc.) do backpropagation for you via automatic differentiation



# How to make deep learning work?

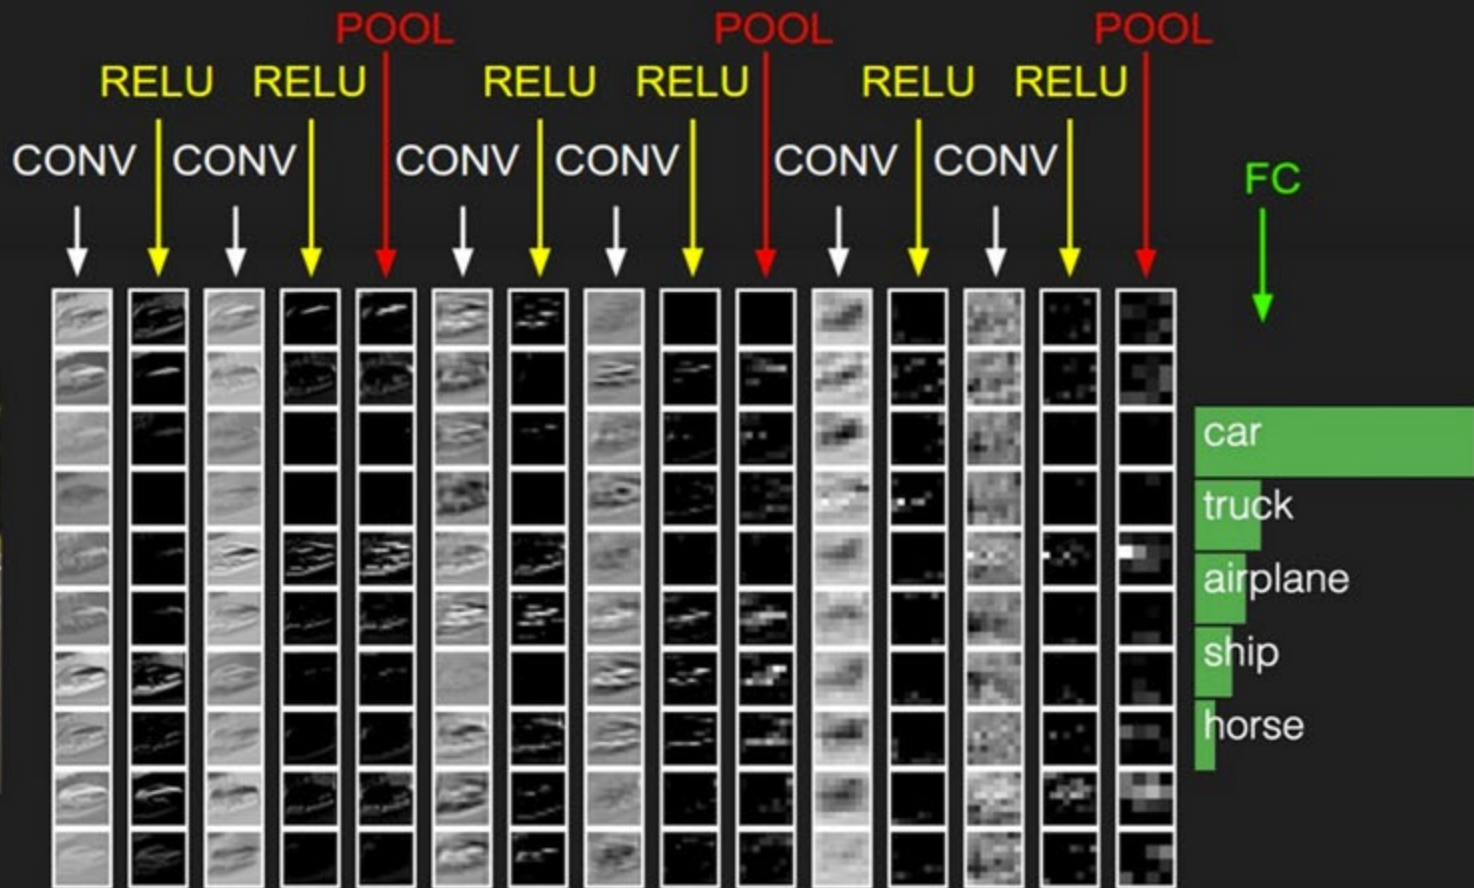
- Carefully designed network **architecture**
- Unsupervised pre-training can be done layer-wise
- Better training algorithms and heuristics such as **good initialization**, **batch normalization**, **weight clip**, **regularization**, etc
- **Drop-out**, **Early Stop** to control complexity, reduce overfitting
- **Transfer learning**, **meta learning**, etc

.

# Structured Neural Models Underlying DL Revolution

- **Convolutional models**
- **Recurrent models**
- **Gated and residual connections**
- **Attention**

# Convolutional Neural Nets




# Vision: Convolutional Models

- For computer vision, a key property that we usually wish to capture is translation invariance
- We would like to have visual “feature detectors” that find something in an image regardless of precisely where it is located
- We do this with a convolutional layer




# Fully Connected Layer

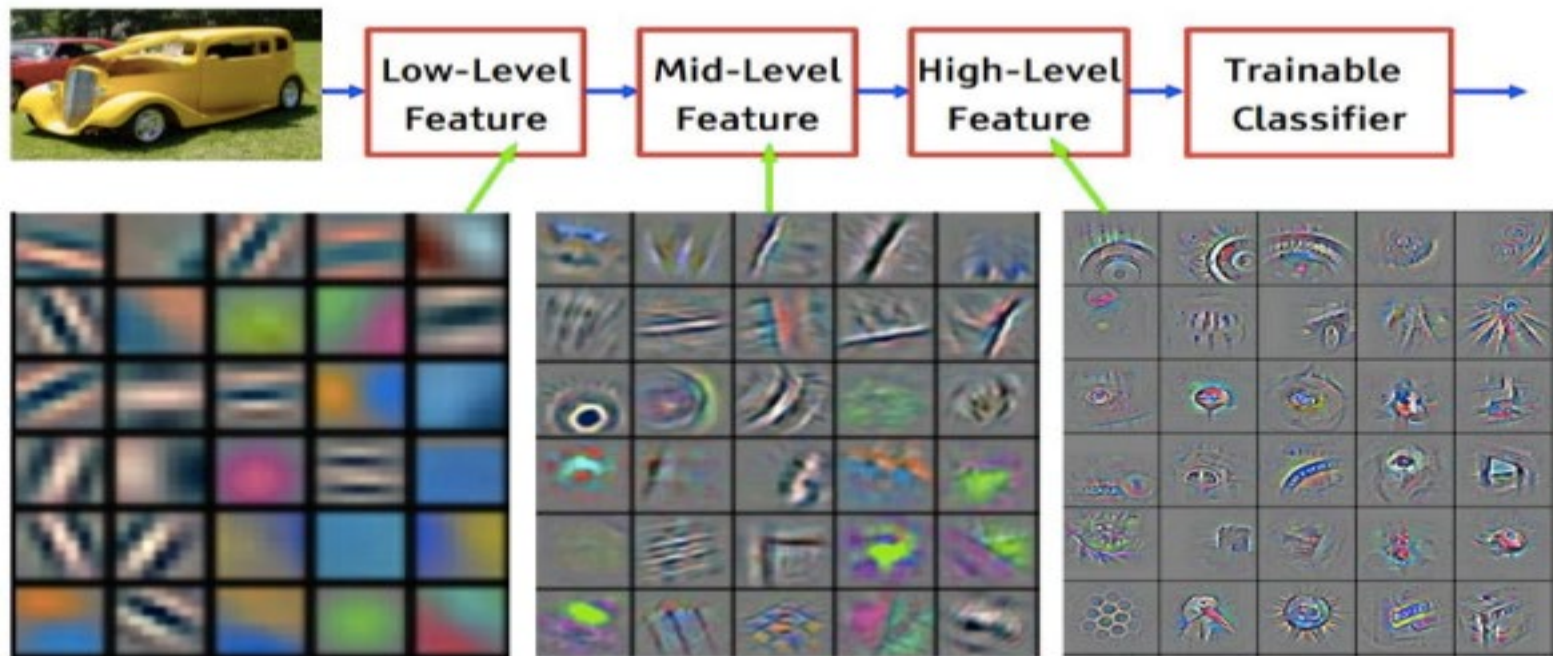
32x32x3 image -> stretch to 3072 x 1



# Convolutional Neural Net



# Features learned by CNN




Feature visualization of convolutional net trained on ImageNet from [Zeiler & Fergus 2013]

# IMAGENET Large Scale Visual Recognition Challenge

Year 2010

NEC-UIUC




[Lin CVPR 2011]

[Lion image](#) by Swissfrog is  
licensed under [CC BY 3.0](#)

Year 2012

SuperVision




[Krizhevsky NIPS 2012]

Figure copyright Alex Krizhevsky, Ilya  
Sutskever, and Geoffrey Hinton, 2012.  
Reproduced with permission.


Year 2014

GoogLeNet



[Szegedy arxiv 2014]

VGG



[Simonyan arxiv 2014]


Year 2015

MSRA



[He ICCV 2015]






This image is licensed under [CC BY-NC-SA 2.0](#); changes made

- Object detection
- Action classification
- Image captioning
- ...



Person

Hammer




# Language: Recurrent Models

- Until now, we've dealt with classifying/generating fixed-size objects.
  - We just resized images to our procrustean bed!
- How can we deal with variable-size inputs, such as the word sequences in human language text or bioinformatic gene sequences?
- We do this with a recurrent layer

# Recurrent Neural Nets (RNN)

**Core idea:** Apply the same weights  $W$  repeatedly



# RNN for Language Modeling


Transition Function  $h_t = f(h_{t-1}, x_t)$

Inputs

- Current word  $x_t \in \{1, 2, \dots, |V|\}$
- Previous state  $h_{t-1} \in \mathbb{R}^d$

Parameters

- Input weight matrix  $W \in \mathbb{R}^{|V| \times d}$
- Transition weight matrix  $U \in \mathbb{R}^{d \times d}$
- Bias vector  $b \in \mathbb{R}^d$





# Building a language model

*Transition Function*  $h_t = f(h_{t-1}, x_t)$


Naïve Transition Function

$$f(h_{t-1}, x_t) = \tanh(W[x_t] + Uh_{t-1} + b)$$

*Element-wise nonlinear transformation*

*Trainable word vector*

*Linear transformation of previous state*



# Building a recurrent language model

*Prediction Function  $p(x_{t+1} = w | x_{\leq t}) = g_w(h_t)$*


Inputs

- i. Current state  $h_t \in \mathbb{R}^d$

Parameters

- i. Softmax matrix  $R \in \mathbb{R}^{|V| \times d}$
- ii. Bias vector  $c \in \mathbb{R}^{|V|}$

*Softmax gives a probability distribution over next words. To generate text, we take the word with max prob., and use it as the input at the next time step*



# Some observations

- **Non-convexity** is no longer an issue. In fact, non-convex models seem to more powerful than convex models.
- We used to worry about bad **local minima**, but it turns out that bad local minima typically don't exist and any different minima seem to be roughly equivalently good. No need to worry!
- Extremely deep models using **residual connections** (ResNets) have powered the progress in computer vision. Need to rethink layered representations?

# Some observations

- **Embedding**

Use of **distributed representations of categorical values** like words allows very effective modeling and sharing of dimensions of similarity



# Some observations

## Rethink overfitting

- Models with orders of magnitude more parameters than there are training data available, if trained with ample amounts of regularization (including new forms like dropout), will outperform simpler models – generalizing better to new data.
- **Practitioner's recipe**: build a sufficiently high-capacity model that it can be trained to 0% training error, and then increase its regularization until it doesn't overfit on dev.

# Some observations

## Architecture matters

- Identify the best architecture for your problem. RNN, CNN, GRU, LSTM ...
- Put in mechanisms to facilitate learning: Dropout, Batch normalization, Attention ...